

Machine Learning Engineer

Lumen Intelligence Labs | Hyderabad, India | Full-time

About the role

We are looking for a Machine Learning Engineer to build, train and deploy models that power our document understanding platform. You will work across the full model lifecycle, from dataset construction to production model deployment.

Key Responsibilities

- Build and train deep learning models for natural language processing tasks
- Design feature engineering pipelines over large document datasets
- Fine-tune and evaluate large language models for extraction and classification
- Deploy models to production and monitor drift, latency and accuracy
- Run rigorous offline evaluation and online A/B experiments
- Collaborate with backend engineers to expose models behind REST APIs

Required Qualifications

- Master's degree in Computer Science, Machine Learning, Statistics or related field
- Minimum 3 years of hands-on machine learning experience
- Strong Python programming skills
- Deep experience with PyTorch or TensorFlow
- Solid grounding in deep learning and neural network architectures
- Practical natural language processing experience
- Experience with pandas, NumPy and scikit-learn
- Strong statistics and experimental design fundamentals
- Experience deploying models to production

Preferred Qualifications

- Experience with large language models and prompt engineering
- Familiarity with MLOps tooling and model registries
- Experience with computer vision
- Exposure to Docker and AWS for model serving
- Publications at machine learning conferences